

Operating Systems: Processes and Threads

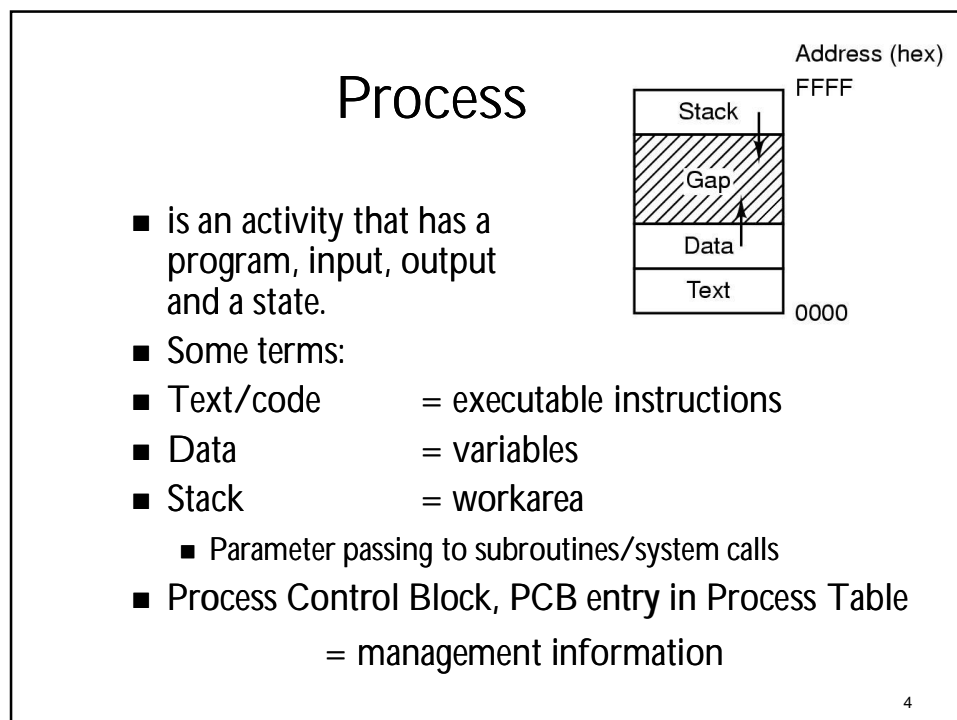
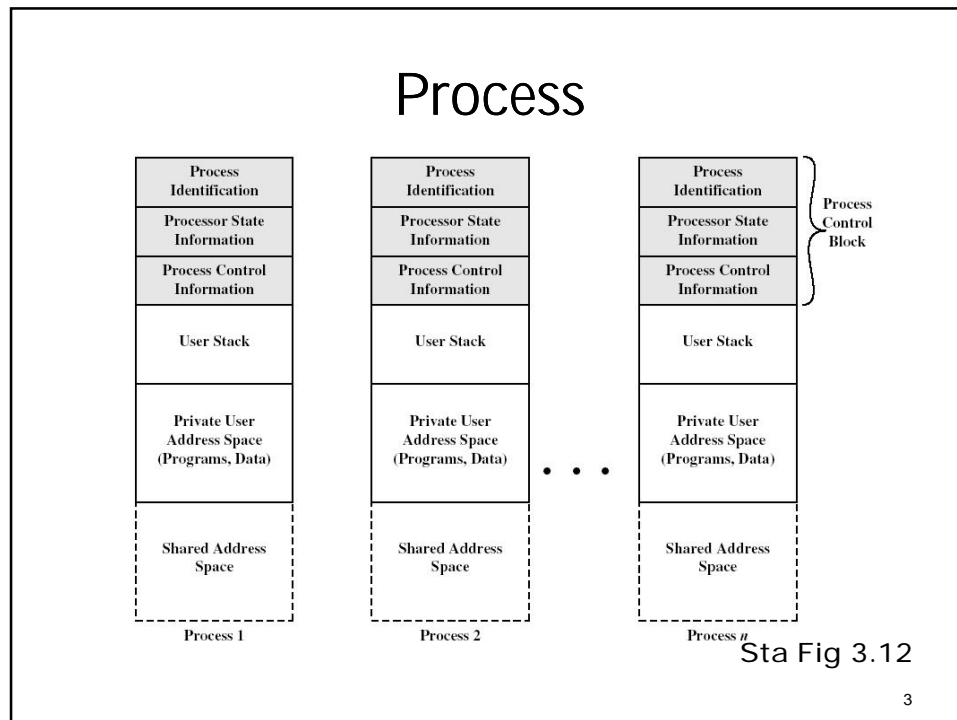
Week 1: Lecture 2, Thu 4.9.2008

Tiina Niklander

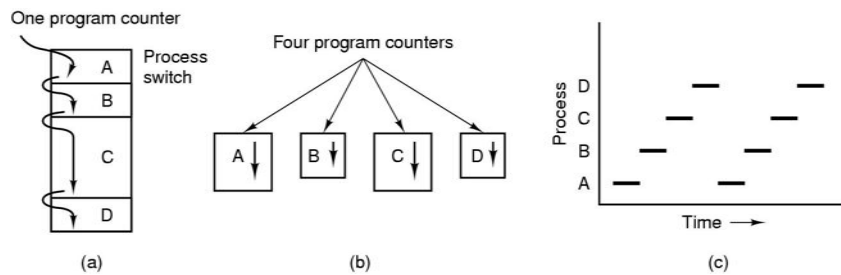
1

Process model

2



Process model



- One physical program counter switches between processes
- Conceptual model of 4 independent, sequential processes
- Only one program active at any instant on one CPU
pseudoparallelism

5

Process Creation

Principal events that cause process creation

1. System initialization
2. Execution of a process creation system
3. User request to create a new process
4. Initiation of a batch job

6

Process Creation

Operating system does

- Create PCB
 - OS 'generates' a unique ID
- Allocate memory for the process
- Initiate PCB
- Link PCB to other structures
 - Place to Ready-queue, link to parent process, etc

7

Process Termination


Conditions which terminate processes

1. Normal exit (voluntary)
2. Error exit (voluntary)
3. Fatal error (involuntary)
4. Killed by another process (involuntary)


8

Process Hierarchies

- Parent creates a child process, child processes can create its own process
- Forms a hierarchy
 - UNIX calls this a "process group"
- Windows has no concept of process hierarchy
 - all processes are created equal



Process States (1)



1. Process blocks for input
2. Scheduler picks another process
3. Scheduler picks this process
4. Input becomes available

- Possible process states
 - running
 - blocked
 - ready
- Transitions between states shown

10

Process states


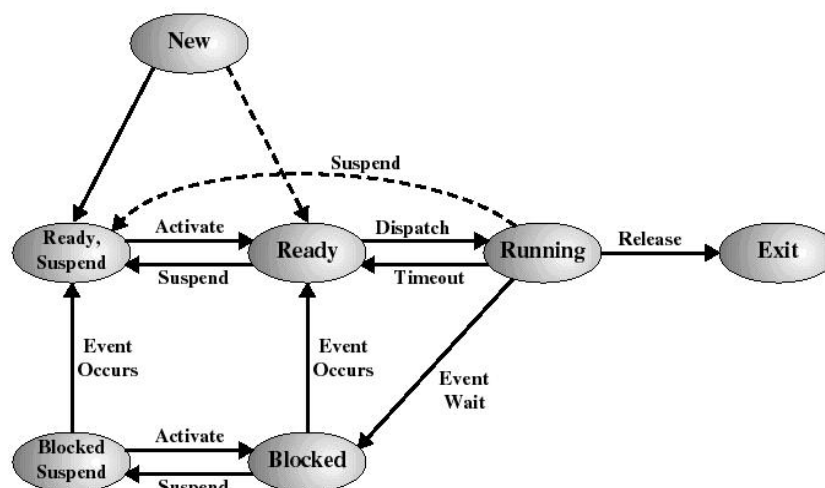


Figure 3.7 Process States for Trace of Figure 3.4

11

Process states



Sta Fig 3.8

12

Process Control Block

Process management	Memory management	File management
Registers Program counter Program status word Stack pointer Process state Priority Scheduling parameters Process ID Parent process Process group Signals Time when process started CPU time used Children's CPU time Time of next alarm	Pointer to text segment Pointer to data segment Pointer to stack segment	Root directory Working directory File descriptors User ID Group ID

Fields of a process table entry

13

Tasks on interrupt

1. Hardware stacks program counter, etc.
2. Hardware loads new program counter from interrupt vector.
3. Assembly language procedure saves registers.
4. Assembly language procedure sets up new stack.
5. C interrupt service runs (typically reads and buffers input).
6. Scheduler decides which process is to run next.
7. C procedure returns to the assembly code.
8. Assembly language procedure starts up new current process.

The process is not always changed in step 6,
if running process can continue.


Process can be changed only at step 6.

14

Threads

15


Word processor with three threads



- One thread interacts with user
- One thread handles reformatting in background
- One thread handles disk backups in background

16

Multithreaded web server using dispatcher



- Dispatcher receives requests
- Handles to an idle worker thread for processing
- Suitable model for such: finite-state machine

```
while (TRUE) {
  get_next_request(&buf);
  handoff_work(&buf);
}
```

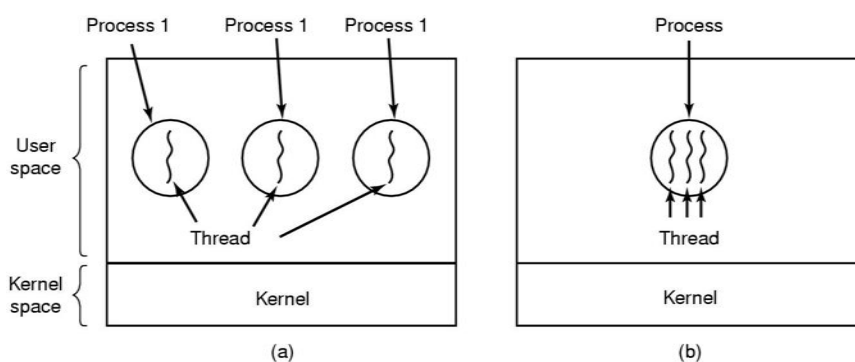
(a)

```
while (TRUE) {
  wait_for_work(&buf)
  look_for_page_in_cache(&buf, &page);
  if (page_not_in_cache(&page))
    read_page_from_disk(&buf, &page);
  return_page(&page);
}
```

(b)

17

Thread model



(a) Three processes each with one thread

(b) One process with three threads

18

Thread Model: Process vs thread

Per process items

Address space
Global variables
Open files
Child processes
Pending alarms
Signals and signal handlers
Accounting information

Per thread items

Program counter
Registers
Stack
State

- Items shared by all threads in a process
- Items private to each thread in thread
- Each thread has its own stack!

19

Why threads?

- Multiple parallel activities. Thread provide easier programming models
- Faster to create and destroy one thread than a whole process
- When one thread waits, another might be able to execute
- Shared data area of threads in one process. Efficient resource sharing and easy communication.

BUT

- Mutual exclusion and synchronization are fully programmer's responsibility (no support from OS or anything else)

⇒ Concurrent
programming
course

20


POSIX threads (pthreads)

IEEE standard for threads for portable thread programs

- `pthread_create()`
 - Create a new thread
- `pthread_exit()`
 - Terminate the calling thread
- `pthread_join()`
 - Wait for a specific thread to exit
- `pthread_yield()`
 - Release the CPU to let another thread run
- Functions for synchronization and mutual exclusion
- And more than 50 other functions


21

User-level vs kernel-level threads



(a) A user-level threads package.

- Kernel (or OS) is not aware of threads, schedules processes
- User process must dispatch threads itself



(b) A threads package managed by the kernel.

- All thread control and dispatching done by the kernel
- No control on the user level

Tan08
Fig 2-16

22

User-level threads

Advantages

- Fast dispatching
 - No mode switch
 - No interrupt
 - No process switch!
- Programmer can freely choose the scheduling mechanism

Disadvantages

- When one thread is blocked on system call, it blocks the whole process (and all other threads)
- Threads of one process cannot be executed on several processors concurrently
 - Remember: kernel dispatches only processes!

23

Kernel-level threads

Advantages


- Threads of one process can be executed simultaneously on multiple processors
- If one thread is Blocked, the other threads of this process may still continue
- Often also the kernel implementation is multithreaded

Disadvantages

- Dispatching a thread has two phases:
 - Interrupt +, interrupt handling and mode switch
 - Dispatcher (and return to user mode)
- Slower than in user-level threads

24

Hybrid Implementations



Multiplexing user-level threads onto kernel-level threads

25

Solaris




Figure 4.15 Solaris Multithreaded Architecture Example

26

Making single-threaded code multithreaded

- Global variables
 - Threads may not be aware of others using the same variable
 - See next slide for solutions
- Library procedures
 - May not be reentrant (second call to procedure before the first one is finished)
 - Solution: rewrite the library or use an excluding jacket
- Signals
 - No simple solutions – difficult already in single thread
- Stack management
 - How to increase a thread's stack in case of stack overflow

27

Threads and global variables

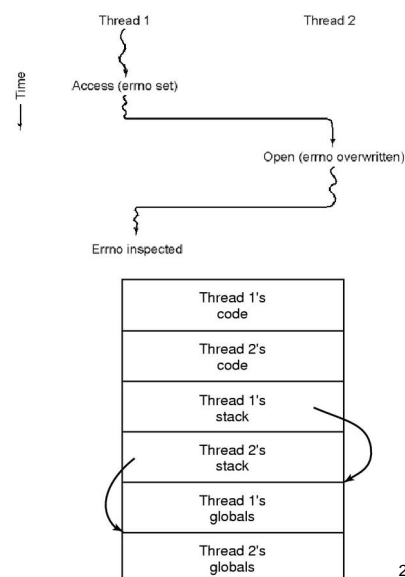
Conflicts between threads over the use of a global variable

One solution:

- Prohibit global variables

Alternative:

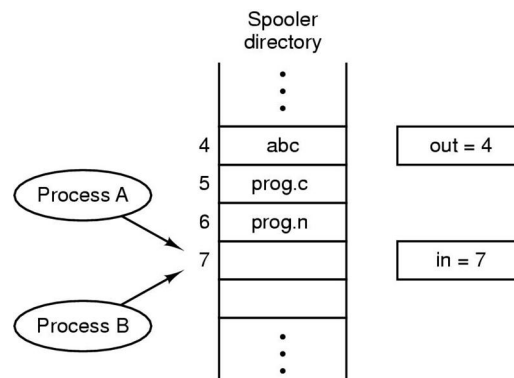
- Private global variables
- Accessing these is tricky, language may not support
- New library procedures



28

Interprocess Communication

Race Conditions



Two processes want to access shared memory at same time

29


Critical Regions (1)

Four conditions to provide mutual exclusion

1. No two processes simultaneously in critical region
2. No assumptions made about speeds or numbers of CPUs
3. No process running outside its critical region may block another process
4. No process must wait forever to enter its critical region

30

Critical Regions (2)



Mutual exclusion using critical regions

31

Mutual Exclusion


- Busy Waiting – occupy the CPU while waiting
- Sleep and Wakeup
 - Semaphores
 - Mutexes
 - Monitors
 - Message Passing
 - Barriers

Concurrent Programming
Rinnakkaisohjelmointi

32

Scheduling

Introduction to Scheduling (1)



- Bursts of CPU usage alternate with periods of I/O wait
 - a CPU-bound process
 - an I/O bound process

33

Introduction to Scheduling (2)

All systems

- Fairness - giving each process a fair share of the CPU
- Policy enforcement - seeing that stated policy is carried out
- Balance - keeping all parts of the system busy

Batch systems

- Throughput - maximize jobs per hour
- Turnaround time - minimize time between submission and termination
- CPU utilization - keep the CPU busy all the time

Interactive systems

- Response time - respond to requests quickly
- Proportionality - meet users' expectations

Real-time systems

- Meeting deadlines - avoid losing data
- Predictability - avoid quality degradation in multimedia systems

Scheduling Algorithm Goals

34

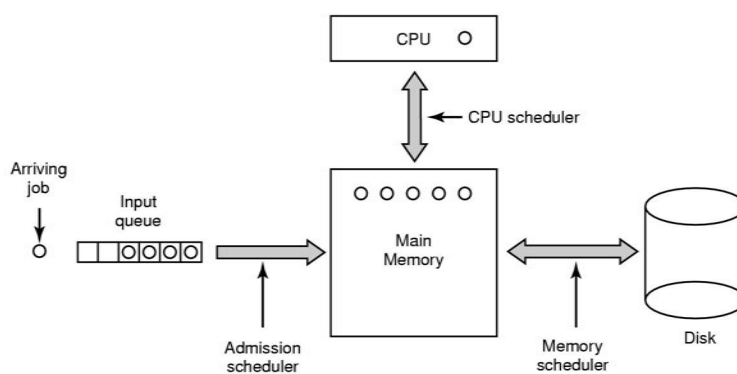
Scheduling in Batch Systems (1)



An example of shortest job first scheduling

35


Scheduling in Batch Systems (2)



Three level scheduling

36

Scheduling in Interactive Systems (1)

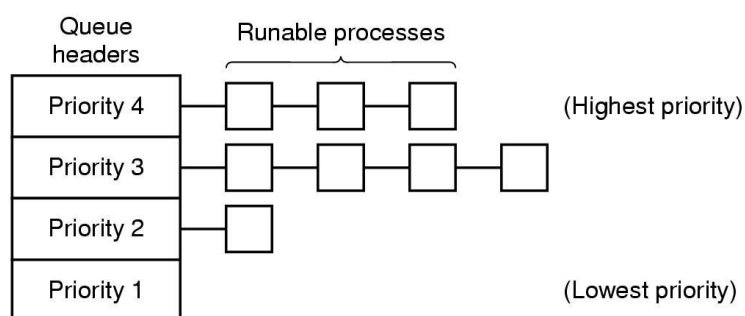


■ Round Robin Scheduling

- list of runnable processes
- list of runnable processes after B uses up its quantum

37

Scheduling in Interactive Systems (2)



A scheduling algorithm with four priority classes

38

Scheduling in Real-Time Systems

Schedulable real-time system

- Given
 - m periodic events
 - event i occurs within period P_i and requires C_i seconds
- Then the load can only be handled if

$$\sum_{i=1}^m \frac{C_i}{P_i} \leq 1$$


39

Policy versus Mechanism

- Separate what is allowed to be done with how it is done
 - a process knows which of its children threads are important and need priority
- Scheduling algorithm parameterized
 - mechanism in the kernel
- Parameters filled in by user processes
 - policy set by user process


40

Thread Scheduling (1)



Possible: A1, A2, A3, A1, A2, A3
 Not possible: A1, B1, A2, B2, A3, B3

User-level



Possible: A1, A2, A3, A1, A2, A3
 Also possible: A1, B1, A2, B2, A3, B3

Kernel-level

Possible scheduling of threads

- 50-msec process quantum
- threads run 5 msec/CPU burst

41